

PROBLEM STATEMENT:- TO PREDICT THE RAIFALL BASED ON VARIOUS FEATURES OF THE DATASET

```
In [1]: import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\Welcome\Downloads\rainfall.csv")
df
```

Out[2]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.9
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.1
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.9
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.8
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.2
...
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527.3
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636.3
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352.7
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592.9
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217.5

641 rows × 19 columns



```
In [3]: df.head()
```

Out[3]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OC
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.9	354.8	326.
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.1	455.6	301.
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.9	454.8	276.
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.8	313.6	167.
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.2	568.0	206.

```
In [4]: df.tail()
```

Out[4]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OC
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527.3	3	
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636.3	2	
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352.7	2	
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592.9	2	
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217.5	1	

```
In [5]: df.isnull().any()
```

```
Out[5]: STATE_UT_NAME    False
DISTRICT                False
JAN                     False
FEB                     False
MAR                     False
APR                     False
MAY                     False
JUN                     False
JUL                     False
AUG                     False
SEP                     False
OCT                     False
NOV                     False
DEC                     False
ANNUAL                  False
Jan-Feb                 False
Mar-May                 False
Jun-Sep                 False
Oct-Dec                 False
dtype: bool
```

```
In [6]: df.fillna(method='ffill',inplace=True)
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: STATE_UT_NAME    0
DISTRICT                0
JAN                     0
FEB                     0
MAR                     0
APR                     0
MAY                     0
JUN                     0
JUL                     0
AUG                     0
SEP                     0
OCT                     0
NOV                     0
DEC                     0
ANNUAL                  0
Jan-Feb                 0
Mar-May                 0
Jun-Sep                 0
Oct-Dec                 0
dtype: int64
```

In [8]: `df.describe()`

Out[8]:

	JAN	FEB	MAR	APR	MAY	JUN	JUL	
count	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	6
mean	18.355070	20.984399	30.034789	45.543214	81.535101	196.007332	326.033697	2
std	21.082806	27.729596	45.451082	71.556279	111.960390	196.556284	221.364643	1
min	0.000000	0.000000	0.000000	0.000000	0.900000	3.800000	11.600000	
25%	6.900000	7.000000	7.000000	5.000000	12.100000	68.800000	206.400000	1
50%	13.300000	12.300000	12.700000	15.100000	33.900000	131.900000	293.700000	2
75%	19.200000	24.100000	33.200000	48.300000	91.900000	226.600000	374.800000	3
max	144.500000	229.600000	367.900000	554.400000	733.700000	1476.200000	1820.900000	15

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 641 entries, 0 to 640
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   STATE_UT_NAME    641 non-null    object
1   DISTRICT         641 non-null    object
2   JAN              641 non-null    float64
3   FEB              641 non-null    float64
4   MAR              641 non-null    float64
5   APR              641 non-null    float64
6   MAY              641 non-null    float64
7   JUN              641 non-null    float64
8   JUL              641 non-null    float64
9   AUG              641 non-null    float64
10  SEP              641 non-null    float64
11  OCT              641 non-null    float64
12  NOV              641 non-null    float64
13  DEC              641 non-null    float64
14  ANNUAL           641 non-null    float64
15  Jan-Feb          641 non-null    float64
16  Mar-May          641 non-null    float64
17  Jun-Sep          641 non-null    float64
18  Oct-Dec          641 non-null    float64
dtypes: float64(17), object(2)
memory usage: 95.3+ KB
```

In [10]: `df.columns`

Out[10]: Index(['STATE_UT_NAME', 'DISTRICT', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May', 'Jun-Sep', 'Oct-Dec'], dtype='object')

```
In [11]: df.shape
```

```
Out[11]: (641, 19)
```

```
In [12]: df['ANNUAL'].value_counts()
```

```
Out[12]: ANNUAL
747.1      9
2080.0     4
1336.5     3
1824.8     3
2814.4     3
..
1037.6     1
907.2      1
944.5      1
1003.3     1
3253.1     1
Name: count, Length: 591, dtype: int64
```

```
In [13]: df['Jan-Feb'].value_counts()
```

```
Out[13]: Jan-Feb
32.7       9
18.2       5
21.4       5
0.8        5
17.5       5
..
107.7      1
87.0       1
101.0      1
135.2      1
65.0       1
Name: count, Length: 399, dtype: int64
```

```
In [14]: df['Mar-May'].value_counts()
```

```
Out[14]: Mar-May
43.5       9
27.9       5
36.6       4
468.6      4
40.4       3
..
16.3       1
23.3       1
49.6       1
20.5       1
232.4      1
Name: count, Length: 511, dtype: int64
```

```
In [15]: df['Jun-Sep'].value_counts()
```

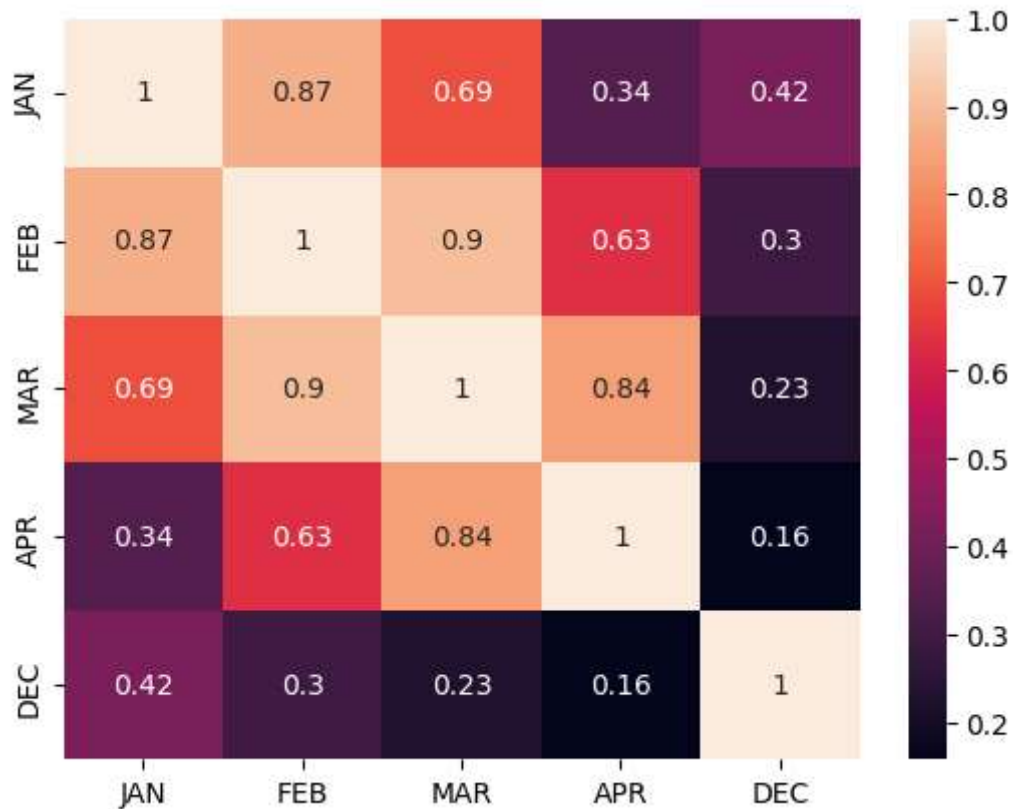
```
Out[15]: Jun-Sep
636.2      9
1386.1      4
385.0       3
1122.3      3
1308.0      3
..
916.9       1
923.5       1
790.3       1
840.7       1
998.5       1
Name: count, Length: 592, dtype: int64
```

```
In [16]: df['Oct-Dec'].value_counts()
```

```
Out[16]: Oct-Dec
34.7        9
174.8        4
49.6         3
27.7         3
183.7        3
..
82.8         1
55.2         1
65.6         1
54.0         1
333.6        1
Name: count, Length: 524, dtype: int64
```

EXPLORATORY DATA ANALYSIS:-

```
In [17]: df=df[['JAN','FEB','MAR','APR','DEC']]
sns.heatmap(df.corr(),annot=True)
plt.show()
```



```
In [18]: df.columns
```

```
Out[18]: Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')
```

```
In [19]: x=df[["FEB"]]
y=df["JAN"]
```

LINEAR REGRESSION:-

```
In [20]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=
```

```
In [21]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.intercept_)
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
coeff_
```

3.6728680241521268

Out[21]:

	coefficient
FEB	0.715365


```
In [22]: score=reg.score(X_test,y_test)
print(score)
```

0.6855837686354153


```
In [23]: predictions=reg.predict(X_test)
```

```
In [24]: plt.scatter(y_test,predictions)
```

Out[24]: <matplotlib.collections.PathCollection at 0x1cd24d58c10>




```
In [25]: df500=df[:][500]
sns.lmplot(x="FEB",y="JAN",order=2,ci=None,data=df500)
plt.show()
```




```
In [26]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
reg.fit(X_train,y_train)
reg.fit(X_test,y_test)
```

Out[26]:

▼ LinearRegression

LinearRegression()

```
In [27]: y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='black')
plt.plot(X_test,y_pred,color='red')
plt.show()
```



```
In [28]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.8131357285061002

RIDGE MODEL

```
In [29]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [30]: features= df.columns[0:5]
target= df.columns[-5]
```

```
In [31]: x=np.array(df['JAN']).reshape(-1,1)
y=np.array(df['FEB']).reshape(-1,1)
```

```
In [32]: x= df[features].values
y= df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)
```

```
In [33]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```


```
In [34]: print("\n Ridge Model:\n")
print("the train score for ridge model is{}".format(train_score_ridge))
print("the test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:

the train score for ridge model is0.9999999792491524
the test score for ridge model is0.9999999887465535

```
In [35]: lr=LinearRegression()
```

```
In [36]: plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=10,color='red')
#plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



LASSO MODEL

```
In [37]: print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is{}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.99912857000705

The test score for ls model is0.9991969731663574

```
In [38]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

Out[38]: <Axes: >




```
In [39]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y_
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.99999999999999198

0.99999999999999254

```
In [40]: plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=10,color='red')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color="green")
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



ELASTIC NET

```
In [44]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
print(regr.score(x,y))
```

```
[ 9.93078025e-01  4.17330119e-03  0.00000000e+00 -2.56844715e-04
  4.33064019e-04]
0.04331617537314614
0.9999905490942288
```

```
In [45]: y_pred_elastic = regr.predict(x_train)
mean_squared_error=np.mean((y_pred_elastic - y_train)**2)
print(mean_squared_error)
```

```
0.004278888506659526
```

CONCLUSION:-

THE SCORE OF LINEAR REGRESSION IS :-0.8131357285061002 THE SCORE OF RIDGE MODEL IS :- 0.9999999887465535 THE SCORE OF LASSO MODEL IS :- 0.9999999999999992 THE SCORE OF ELASTIC NET IS :-0.004278888506659526 AMONG ALL MODELS LASSO YEILD HIGHEST ACCURACY.SO,WE PREFER LASSO MODEL FOR THIS DATA SET